LOCAL LLMS · YOUR HARDWARE · YOUR DATA

4 000 000 Files

40 Knobs

1 Right Answer.

LLAMATOASTER — THE APPLIANCE THAT FINDS IT.

4 000 000 files. 40 knobs. One right answer.

Local inference shouldn't require becoming a llama.cpp expert.

The Problem

01 **One bad knob — silent quality loss, or a hard crash**

Wrong KV cache → degraded tokens. Context too large → OOM.

02 **One good config on GPU X — broken on GPU Y**

Reddit recipes fail across architectures, drivers, OS versions.

03 **The stale oracle**

Ask ChatGPT “best local LLM 2026” → confident, outdated folklore.

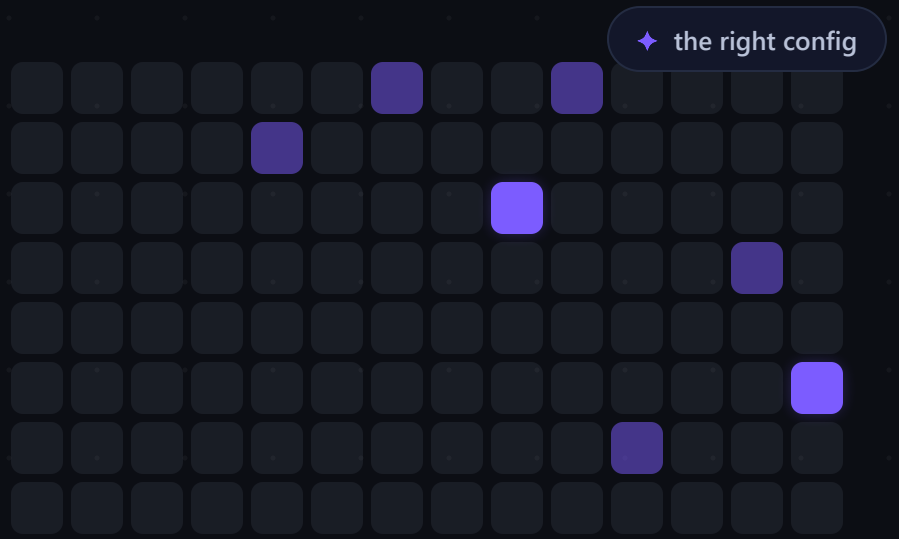

~200 000 GGUF ready

≈4 000 000 files

10+ architectures

The Missing Half

- LMArena · Open LLM Leaderboard · MMLU · GPQA · HumanEval · **MT-Bench** — every release is scored nightly.
- But almost nothing answers the other half: **tok/s · latency · genuinely usable context — on YOUR hardware.**

Millions of files and mountains of quality scores — while performance is still folklore and trial-and-error.

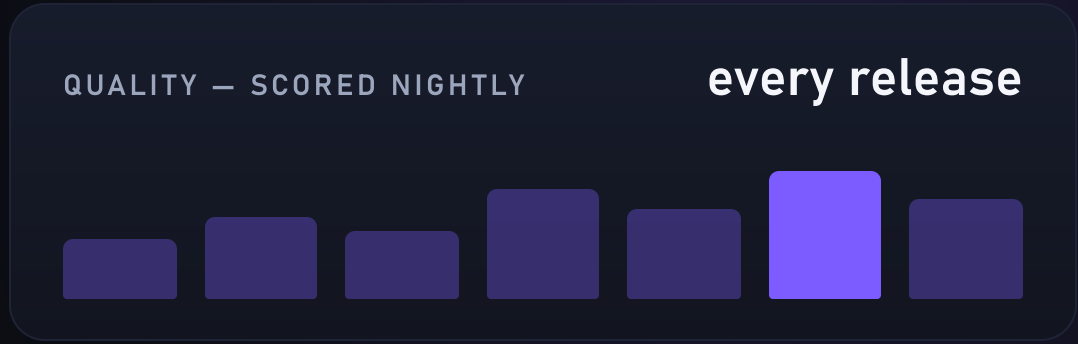

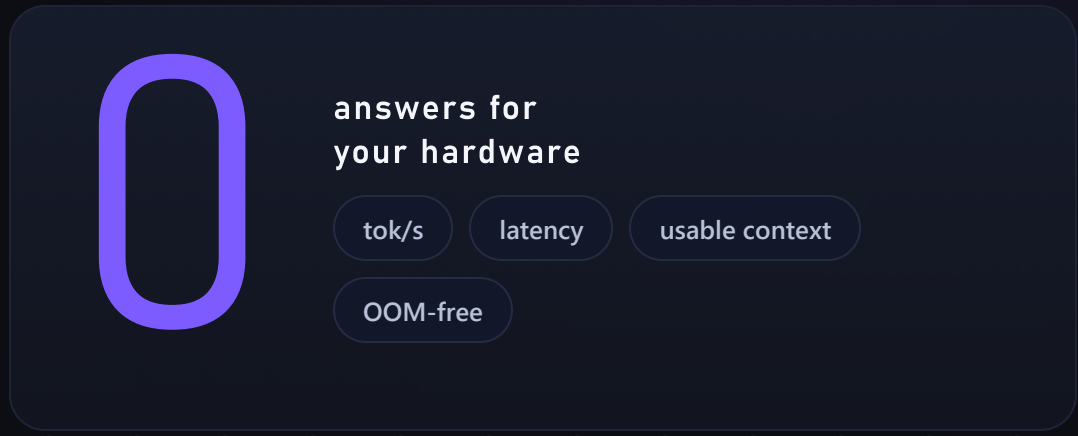

Why Now

Qwen3.8 27B ≈ previous-generation flagship quality — on an ordinary PC.

- ✕ **Blindly chasing the biggest model**
  → swap, crash, 0.3 tok/s
- ✕ **Chasing hype**
  → Llama 4, when Qwen3.8 is better for YOUR task
- ✕ **Following last month's advice**
  → already stale

Proprietary APIs: outages, refusals, price changes — nothing you control. Your data stays on your machine.

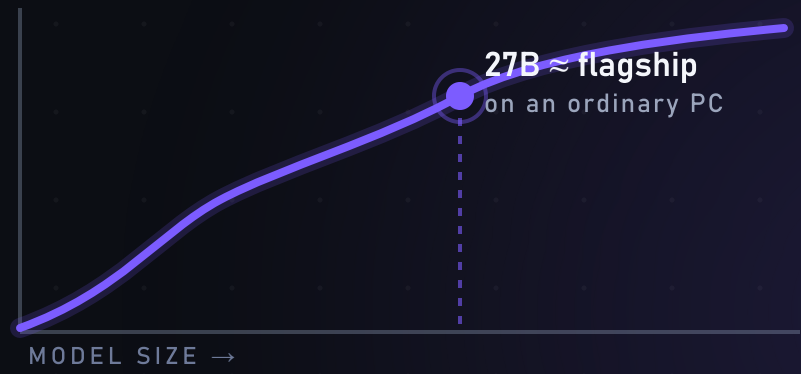

Gemma 3 · Kimi K2 · GLM-4 · **every month** a new “impossible” small model

The Fix

One command connects any GPU or CPU box — **pull-only, zero open ports**, no firewall gymnastics.

- You state intent: **goal × workload × target context**.
- The grid builds itself. **The sweep runs itself.**

5 min set up tonight

Intent in · **Decision out**

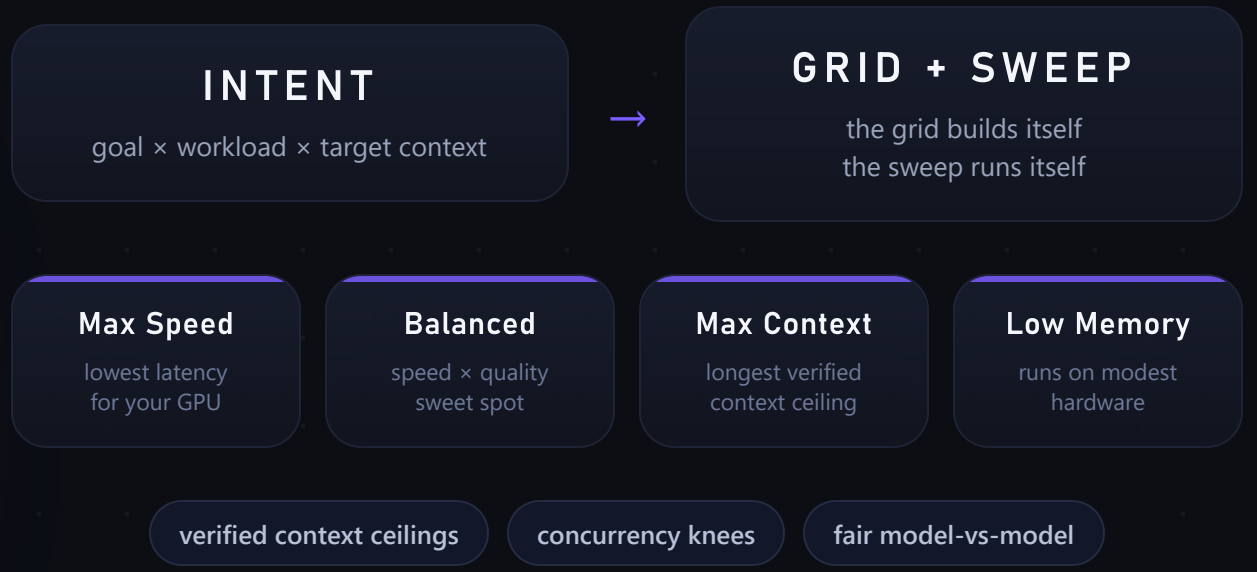

Trust the Card

- ✓ **Eligibility gates.** stability, suspect samples, stddev floors — timer-bug results like **1e6 tok/s** go to the trash, not the average.
- ✓ **Memory from real placement.** per-tensor GGUF placement catches "31/31 offloaded" while secretly running in system RAM.
- ✓ **Fair by construction.** a different machine, build, or GPU rejects the member, not the verdict.
- ✓ **Tamper-evident.** a hash per row, methodology versions never mixed.

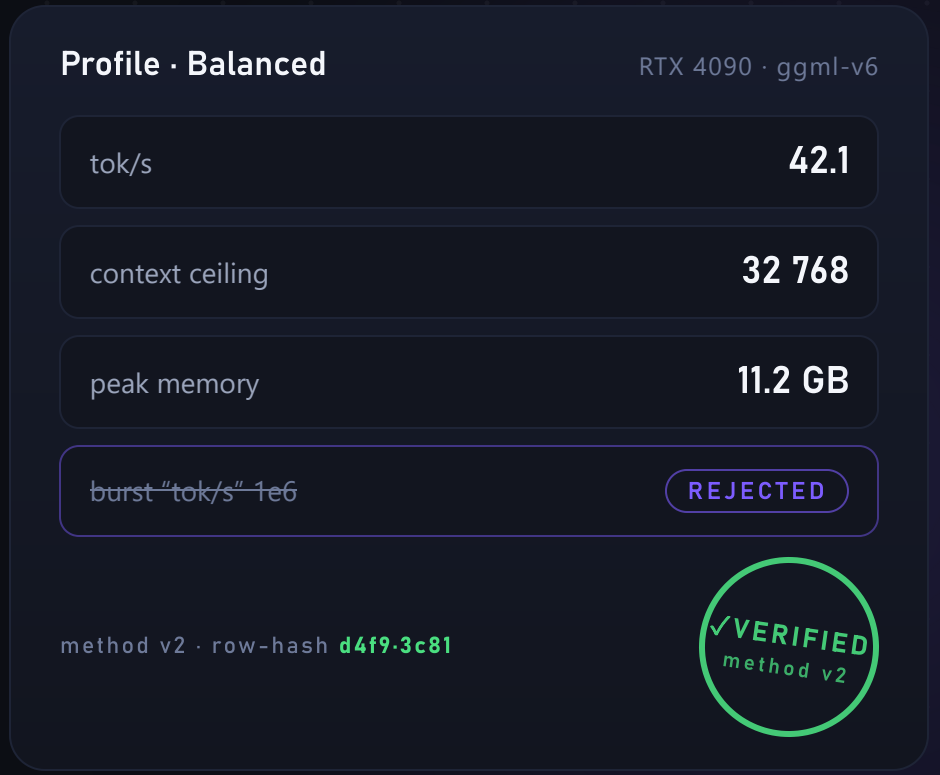

The Competitors

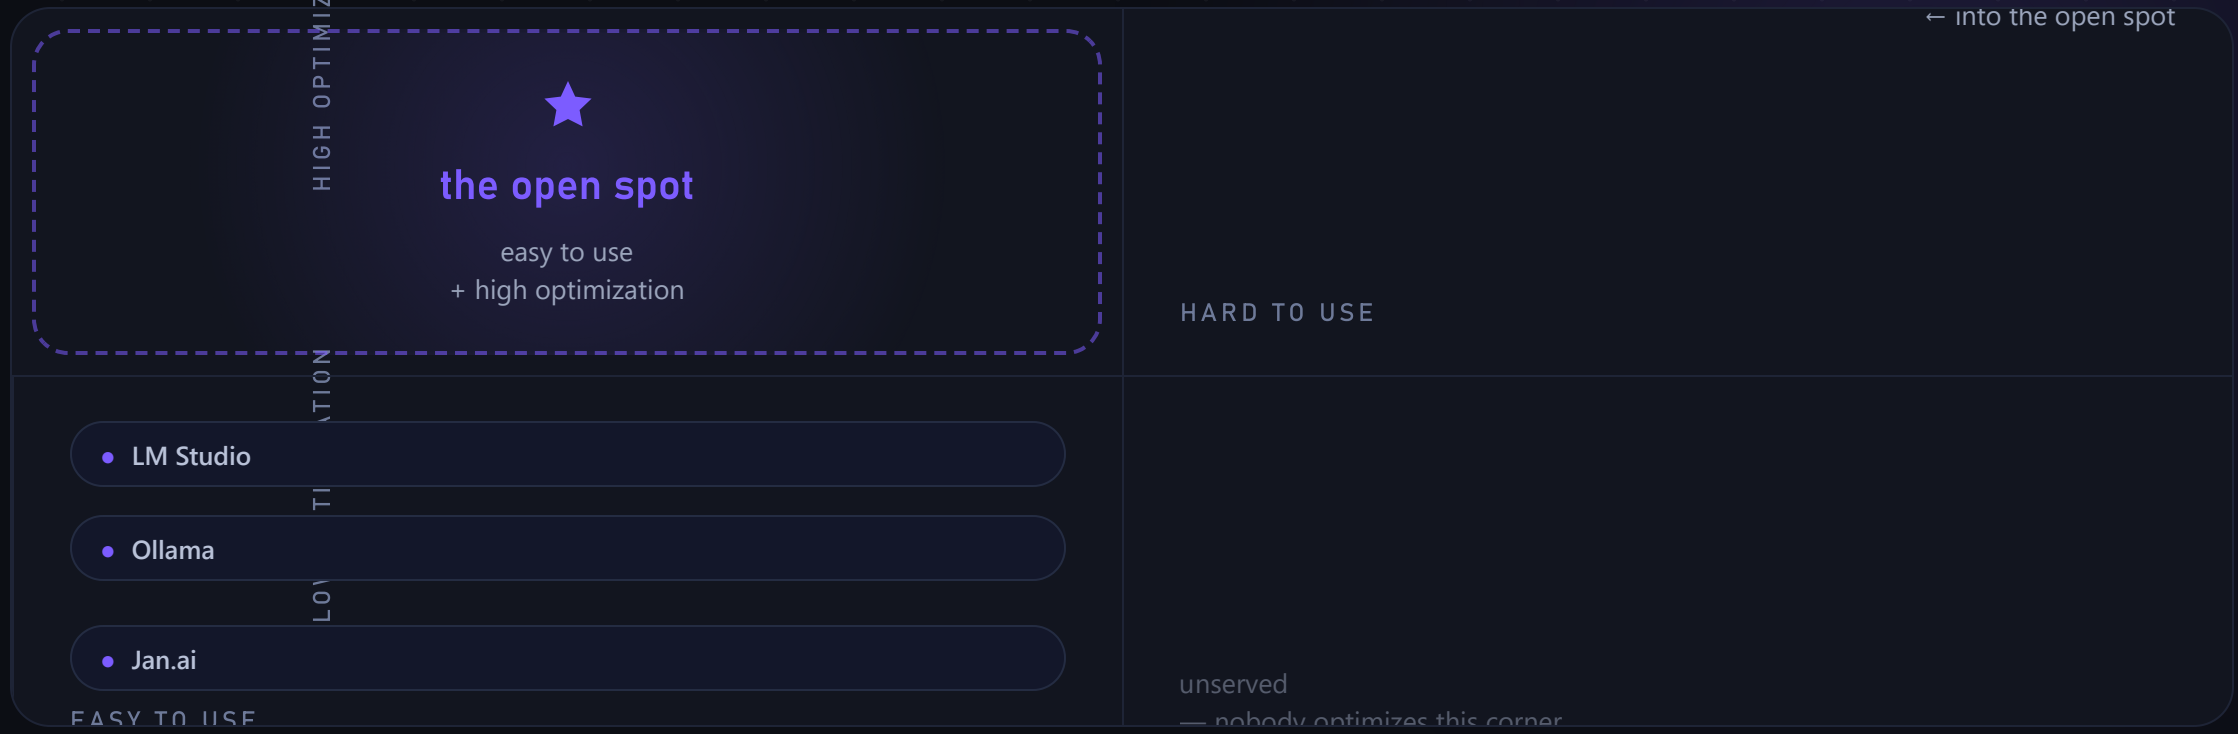

WHAT THEY DO

- ✓ pick a model that “probably works”
- ✓ abstract complexity · one-click install

WHAT THEY DON'T DO

- ✗ **find optimal config** for YOUR hardware
- ✗ **sweep** parameter combinations
- ✗ detect **silent quality degradation**
- ✗ community grid · **tamper-evident** benchmark cards

The Ask

1 command
connects your machine

llamatoaster.com/ **benchmark**

01 / 03
Run your rig tonight
no port forwarding, no manual build — llama.cpp installs from the web UI

02 / 03
Join the community grid
opt into the k-anonymized set → the whole catalog becomes searchable for everyone

03 / 03
Star it while it's early
llamatoaster.com/benchmark — one command connects your machine