

LOCAL LLMS · YOUR HARDWARE · YOUR DATA

≈ 4 000 000 · 14 · 1

GGUF FILES

SWEPT AXES

ANSWER

LLAMATOASTER

the appliance that finds it.

Running a model locally shouldn't require becoming a llama.cpp expert.

# Nobody can read a three-million-model index

~3 000 000

models on Hugging Face

~200 000


of them GGUF — already shaped for llama.cpp

≈4 000 000

files once every quantization is counted

10+ Llama · Qwen · Gemma · Mistral · Phi · DeepSeek · GLM · Kimi · MoE · multimodal

♦ the right one for your box



The index grows weekly. Last month's "best" is already stale.  
So everyone picks by folklore: a Reddit thread said that one is good.

# Choosing is one haystack. Configuring it is another.

14

**axes, cross-multiplied**

on every run, on every machine, for every model



**The first config that doesn't crash wins — not the best one.**  
**Complexity isn't overcome; it filters people out.**

01

**One bad knob — silent quality loss, or a hard crash**

Wrong KV cache → degraded tokens. Context too large → OOM at hour three.

02

**One good config on GPU X — broken on GPU Y**

Copied recipes fail across architectures, drivers, OS versions.

03

**The stale oracle**

Ask a chatbot for the best local setup and you get confident folklore from a training set that closed before the model you want existed.

# Quality is measured to death. Performance is a black box.

- Dozens of leaderboards rank "smartness" — every release scored within days.  
LMArena · Open LLM Leaderboard · MMLU · GPQA · HumanEval · MT-Bench.
- Almost nothing answers the other half:  
**tok/s · time-to-first-token · genuinely usable context — on YOUR hardware.**

Mountains of quality scores. Performance is still trial and error.  
The hardware was almost always fine. The choice layer is what fails.

QUALITY — SCORED NIGHTLY

every release



answers for  
your hardware

tok/s

TTFT

usable context

OOM-free

# Models crossed the line — but accessible ≠ obvious

Proprietary APIs bring outages, refusals, price changes — none of it yours.  
Your data stays on your machine, off every provider's logs.

- ✗ **Chasing the biggest model**  
→ swap, crash, 0.3 tok/s
- ✗ **Chasing hype**  
→ the loud model instead of the right one for YOUR task
- ✗ **Following last month's advice**  
→ already stale

## THE LINE MODELS JUST CROSSED

**A 27B–32B quantized model on one consumer GPU now lands where last generation's flagship did.**

Every month brings another "impossible" small model.

The missing piece is maximum effect for your workload — not maximum parameters, not the loudest release.

no outages

no refusals

no price changes

✦ **nothing leaves your box**

# An orchestrator that answers, not just measures

One command connects any GPU or CPU box. Pull-only: the worker long-polls out, the server never dials in. Nothing to port-forward, nothing to firewall-open.

**You state intent — goal × workload × target context.**

The grid builds itself. The sweep runs itself, in three linked stages: **tuning** → **refine** → **sweep**.

**Intent in.**

**Decision out.**

Set up tonight. By morning: "run this exact configuration" — not "go figure it out."



Out come four profile cards — Max Speed · Balanced · Max Context · Low Memory — each with the exact command line and a human-readable reason it won.

```
llama-bench -m model.gguf -fa on -ctk q8_0 -ctv q8_0 -ngl 31 -b 512 -ub 512 -d 4096
```

— ALREADY BUILT

# Not a roadmap — this is what runs today



## Fleet

many machines, one queue; device-flow enrolment, per-user isolation in SQL



## Models

search Hugging Face, download a quant straight onto a chosen machine



## Builds

install, activate and delete llama.cpp releases from the web UI, no compiler



## Probe ladder

6 modes bisect the real context ceiling instead of guessing it



## Curves

tok/s and TTFT against growing context; the concurrency knee, derived on read



## Speculative decoding

MTP runs benchmarked through llama-server, not just llama-bench



## Quality

llama-perplexity PPL and KLD against a pinned corpus, so degradation is measured



## Assistant

your own hardware, models and results as context, plus opt-in community aggregates



## Exchange

JSON · CSV · Markdown export, and importable bundles that keep their methods

Nine surfaces, one queue, **one command to attach a machine**.  
Every one of these is in the repository today — none of it is a promise.

**RUNS TODAY**

# Every number carries its methods

- ✓ Four eligibility gates — stability, suspect samples, missing pairs, caveat flags — send timer-bug results like 1e6 tok/s to the tally, not into the average. Every rejection is counted and shown: a silent zero-card outcome would be indistinguishable from a bug.
- ✓ Memory is estimated per-tensor from the real GGUF, not from file size ÷ layers — it catches the config that logs "31/31 layers offloaded" while quietly paging into system RAM.
- ✓ No cell renders a number without naming its source: measured here, derived from llama-bench, or unavailable. Method versions are never averaged across.
- ✓ Comparisons are fair by construction — a different machine, build or GPU rejects the member, not the verdict. Exports carry a hash per row, and community aggregates are k-anonymised in SQL: never a group under 5 contributors, opt-in, never your own rows.

## Max Speed

method v3 · 5 repeats

tg q8\_0/q8\_0 ngl 31

42.7 tok/s

pp q8\_0/q8\_0 ngl 31

918 tok/s

tg f16/f16 ngl 33

1.0e6 tok/s

rejected · suspect samples

config\_hash a91f2c7e... · ctx 16 384 verified

VERIFIED  
PROBED

# Everyone else solves "it runs." We solve "run it this way."

| TOOL                     | WHAT IT GIVES YOU                        | WHAT IT NEVER ANSWERS                         |
|--------------------------|------------------------------------------|-----------------------------------------------|
| Ollama · LM Studio · Jan | One-click run, sane defaults             | Is this the best config for THIS box?         |
| llama-bench              | Honest numbers for one command line      | Which of 4 000 combinations should I run?     |
| Leaderboards             | Which model is smartest                  | What will it do on MY GPU?                    |
| Ask a chatbot            | An instant confident answer              | Whether it was ever true                      |
| LlamaToaster             | A ranked decision, with the command line | Output quality — until you run a quality pass |

The gap isn't running models. **It's choosing how.**

# Run your rig tonight

## 1 command connects your rig

a spare GPU box, an old laptop, or the CPU on your VPS

[github.com/ noname9006/LlamaToaster](https://github.com/noname9006/LlamaToaster)

[llamatoaster.com](https://llamatoaster.com)

01

### Attach a machine

One command. Pull-only — nothing to port-forward, no manual llama.cpp build.

02

### Share your numbers

Opt in and your results join the k-anonymised community set — never under 5 contributors.

03

### Star it while it's early

Solo GPU owners, mixed-hardware teams, and anyone publishing reproducible numbers.